

Chapter 1 – The Big Picture

Main point: Introduces Linux as layers of abstraction: hardware at the bottom, the kernel in the middle, and user space (applications, shells, tools) on top. It highlights the kernel’s key jobs: process management, memory management, device drivers, and system calls, plus the concept of users and permissions.

Why it matters: This gives you a mental map of the system. When something breaks or behaves oddly (slow program, “no space,” weird device issues), you can pinpoint which *layer* is involved and avoid random trial-and-error.

Chapter 2 – Basic Commands and the Directory Hierarchy

Main point: Covers everyday shell usage (Bourne shell /bin/sh and relatives), basic file and directory commands (ls, cp, mv, rm, cd, mkdir, grep, less, find, etc.), environment variables, processes, permissions, archiving, and the standard Linux directory layout (/ , /usr, /var, /home, etc.).

Why it matters: These are the “vocabulary” of Linux. Understanding *why* directories are organized this way and how permissions and processes work prepares you to understand later chapters that dive into the implementation details of filesystems, daemons, and the boot process.

Chapter 3 – Devices

Main point: Explains Linux device files under /dev, how they represent hardware, and how subsystems like sysfs and udev map physical devices to these files. It surveys typical device naming schemes: disks (/dev/sd*, /dev/nvme*), terminals (/dev/tty*), serial ports, audio devices, and more.

Why it matters: When dealing with disks, USB sticks, serial gadgets, or audio hardware, Linux exposes them as files. Understanding this abstraction is crucial for tasks like imaging disks, diagnosing device problems, or writing scripts that interact with hardware reliably.

Chapter 4 – Disks and Filesystems

Main point: Describes how Linux sees block devices, how disks are partitioned, and how filesystems (ext4, XFS, etc.) are created, mounted, and managed. It covers partition tables, filesystem types, and how data is structured and stored.

Why it matters: This is the foundation of storage administration: installing Linux, resizing disks, setting up new volumes, or recovering from disk issues all rely on understanding partitions, filesystems, and mount points.

Chapter 5 – How the Linux Kernel Boots

Main point: Walks through the boot sequence: firmware (BIOS/UEFI), bootloader (like GRUB), loading the kernel, passing parameters, and the kernel’s early initialization (mounting the root filesystem, starting the first process).

Why it matters: Boot failures are among the most intimidating problems. Knowing the boot chain helps you debug “kernel panic,” missing init, wrong root device, or misconfigured bootloader issues instead of reinstalling blindly.

Chapter 6 – How User Space Starts

Main point: After the kernel is up, “user space” begins: the init system (traditionally SysV init, now usually systemd), service management, runlevels/targets, and how background daemons are spawned and supervised.

Why it matters: This explains *how services actually get started*—web servers, databases, logging daemons, etc. If a service won’t start or restarts constantly, understanding the init system is how you fix it systematically.

Chapter 7 – System Configuration: Logging, Time, Batch Jobs, Users

Main point: Covers system-wide configuration topics: logging (syslog/journal), system time and time synchronization, scheduled jobs (cron, at), and user and group management.

Why it matters: These are core admin tasks: tracking what happens on a system (logs), ensuring correct timestamps (important for security and distributed systems), automating routine tasks (backups, clean-ups), and controlling who can do what on the machine.

Chapter 8 – Processes and Resource Utilization

Main point: Takes a deeper look at processes, threads, signals, and process relationships, plus resource usage (CPU, memory, I/O) and tools like ps, top, htop, and related utilities.

Why it matters: Performance troubleshooting is mostly about understanding how processes use resources. This chapter gives you the conceptual tools and commands to investigate slowdowns, runaway processes, and resource limits.

Chapter 9 – Understanding Your Network and Its Configuration

Main point: Explains network fundamentals from a Linux perspective: interfaces, IP addressing, routing, DNS, and how network configuration is stored and applied on modern distributions.

Why it matters: Nearly every Linux system is networked. Understanding how IP addresses, routes, and DNS are configured is essential for debugging connectivity issues, setting up servers, and securing services.

Chapter 10 – Network Applications and Services

Main point: Examines specific network services and applications (like SSH, web servers, mail, etc.), how they’re configured, and how they interact with the network stack.

Why it matters: This is where “Linux internals” meet real-world services. If you run or maintain servers, you need to know how these applications use ports, configuration files, and authentication mechanisms.

Chapter 11 – Introduction to Shell Scripts

Main point: Introduces shell scripting as a way to automate tasks, glue together commands, and create small tools using the shell’s control structures (variables, conditionals, loops, functions, etc.).

Why it matters: Knowing how Linux works is powerful; scripting lets you *operationalize* that knowledge. You can write your own maintenance scripts, deployment helpers, or monitoring checks instead of doing everything manually.

Chapter 12 – Network File Transfer and Sharing

Main point: Covers tools and protocols for moving and sharing files over networks: SCP/SFTP, rsync, NFS, Samba, and related technologies.

Why it matters: Almost every real system needs to move data around—backups, deployments, shared storage. Understanding how network file transfer and sharing work lets you build efficient, secure workflows for data.

Chapter 13 – User Environments

Main point: Discusses how user environments are constructed: shells, dotfiles, environment variables, profiles, and customization of the command-line experience.

Why it matters: A powerful Linux user environment can massively boost productivity. Knowing how login shells load configuration and how to structure dotfiles helps you create a consistent, reproducible setup across machines.

Chapter 14 – Linux Desktop and Printing

Main point: Provides a survey of the Linux graphical desktop (X/Wayland, desktop environments, display managers) and how printing works under Linux.

Why it matters: Not all Linux use is server-side; many users run graphical desktops. Understanding how the GUI stack and printing subsystem fit into the rest of the system helps you troubleshoot graphics, user sessions, and printer issues.

Chapter 15 – Development Tools

Main point: Introduces common development tools on Linux: compilers, debuggers, build systems, and related utilities used to build and test software.

Why it matters: Linux is a primary platform for software development. Knowing the tools available and how they integrate into the system helps both developers and sysadmins support build and CI environments.

Chapter 16 – Compiling Software from C Source Code

Main point: Goes through the full pipeline for compiling C programs: preprocessing, compilation, linking, libraries, and how build systems and makefiles orchestrate these steps.

Why it matters: Many core Linux components and tools are written in C. Understanding compilation and linking helps you build software from source, debug compilation errors, and reason about binary compatibility and libraries.

Chapter 17 – Virtualization

Main point: Discusses virtualization and related techniques (virtual machines, containers, etc.) and how Linux supports them through features like namespaces and control groups.

Why it matters: Modern infrastructure heavily depends on virtualization and containers. Understanding how Linux isolates and manages virtualized workloads is essential for cloud, DevOps, and modern server administration.